

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
import numpy as np
import os
```

```
from keras.models import Sequential
from keras.layers import Dense, Conv2D, Dropout, Flatten
from keras.constraints import maxnorm
from tensorflow.keras.optimizers import Adam
```

```
from keras.layers.convolutional import Convolution2D, MaxPooling2D
from keras.callbacks import ModelCheckpoint, LearningRateScheduler
from keras.callbacks import ReduceLROnPlateau, EarlyStopping
```

```
from keras.utils import np_utils
import matplotlib.pyplot as plt
from keras.preprocessing.image import ImageDataGenerator
```

```
x_train = '/content/drive/MyDrive/tropical_fruits/train'
x_test = '/content/drive/MyDrive/tropical_fruits/test'
```

```
x_train = ImageDataGenerator(rescale=1/255)
x_test = ImageDataGenerator(rescale=1/255)
```

```
x_train_data = x_train.flow_from_directory(
    directory= r"/content/drive/MyDrive/tropical_fruits/train",
    target_size=(224,224),
    batch_size=3,
    class_mode='categorical'
)
x_test_data = x_test.flow_from_directory(
    directory= r"/content/drive/MyDrive/tropical_fruits/test",
    target_size=(224,224),
    batch_size=3,
    class_mode='categorical'
)
```

```
Found 66 images belonging to 10 classes.
Found 28 images belonging to 10 classes.
```

```
x_train_data.class_indices
```

```
{'apple': 0,
 'banana': 1,
 'cherry': 2,
 'coconut': 3,
 'durian': 4,
 'kiwi': 5,
 'mango': 6,
 'orange': 7,
 'pomelo': 8,
 'water melon': 9}
```

```
model = Sequential()

model.add(Conv2D(32,(3,3),input_shape=(224,224,3),padding='same',activation='relu'))
model.add(Dropout(0.2))

model.add(Conv2D(32,(3,3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2)))

model.add(Conv2D(64,(3,3),activation='relu',padding='same'))
model.add(Dropout(0.2))

model.add(Conv2D(64,(3,3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2)))

model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(Dropout(0.2))

model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2)))

model.add(Flatten())
model.add(Dropout(0.2))

model.add(Dense(1024,activation='relu'))
model.add(Dropout(0.2))

model.add(Dense(512,activation='relu'))
model.add(Dropout(0.2))

model.add(Dense(10,activation='softmax'))
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 224, 224, 32)	896
dropout (Dropout)	(None, 224, 224, 32)	0
conv2d_1 (Conv2D)	(None, 224, 224, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 112, 112, 32)	0
conv2d_2 (Conv2D)	(None, 112, 112, 64)	18496
dropout_1 (Dropout)	(None, 112, 112, 64)	0
conv2d_3 (Conv2D)	(None, 112, 112, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 64)	0
conv2d_4 (Conv2D)	(None, 56, 56, 128)	73856
dropout_2 (Dropout)	(None, 56, 56, 128)	0
conv2d_5 (Conv2D)	(None, 56, 56, 128)	147584
max_pooling2d_2 (MaxPooling2D)	(None, 28, 28, 128)	0

flatten (Flatten)	(None, 100352)	0
dropout_3 (Dropout)	(None, 100352)	0
dense (Dense)	(None, 1024)	102761472
dropout_4 (Dropout)	(None, 1024)	0
dense_1 (Dense)	(None, 512)	524800
dropout_5 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 10)	5130

```

=====
Total params: 103,578,410
Trainable params: 103,578,410
Non-trainable params: 0

```

```
from tensorflow.keras.optimizers import SGD
```

```
model.compile(optimizer=Adam(learning_rate=0.0005),
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

```
history = model.fit(x_train_data, epochs= 30, batch_size= 32, verbose= 1,
                   validation_data= x_test_data)
```

```

Epoch 1/30
22/22 [=====] - 3s 85ms/step - loss: 2.2984 - accuracy: 0.0303 - val
Epoch 2/30
22/22 [=====] - 2s 77ms/step - loss: 2.2971 - accuracy: 0.0909 - val
Epoch 3/30
22/22 [=====] - 2s 77ms/step - loss: 2.2982 - accuracy: 0.0758 - val
Epoch 4/30
22/22 [=====] - 2s 78ms/step - loss: 2.2986 - accuracy: 0.1515 - val
Epoch 5/30
22/22 [=====] - 2s 75ms/step - loss: 2.2957 - accuracy: 0.0758 - val
Epoch 6/30
22/22 [=====] - 2s 78ms/step - loss: 2.2972 - accuracy: 0.0606 - val
Epoch 7/30
22/22 [=====] - 2s 76ms/step - loss: 2.2966 - accuracy: 0.0909 - val
Epoch 8/30
22/22 [=====] - 2s 76ms/step - loss: 2.2995 - accuracy: 0.1061 - val
Epoch 9/30
22/22 [=====] - 2s 76ms/step - loss: 2.2993 - accuracy: 0.1061 - val
Epoch 10/30
22/22 [=====] - 2s 71ms/step - loss: 2.2967 - accuracy: 0.1515 - val
Epoch 11/30
22/22 [=====] - 2s 78ms/step - loss: 2.2984 - accuracy: 0.1364 - val
Epoch 12/30
22/22 [=====] - 2s 76ms/step - loss: 2.2989 - accuracy: 0.0909 - val
Epoch 13/30
22/22 [=====] - 2s 79ms/step - loss: 2.2977 - accuracy: 0.0909 - val
Epoch 14/30
22/22 [=====] - 2s 79ms/step - loss: 2.2992 - accuracy: 0.0606 - val
Epoch 15/30
22/22 [=====] - 2s 76ms/step - loss: 2.2978 - accuracy: 0.0909 - val
Epoch 16/30
22/22 [=====] - 2s 79ms/step - loss: 2.2974 - accuracy: 0.0606 - val
Epoch 17/30
22/22 [=====] - 2s 77ms/step - loss: 2.2990 - accuracy: 0.0909 - val
Epoch 18/30
22/22 [=====] - 2s 79ms/step - loss: 2.2982 - accuracy: 0.0909 - val
Epoch 19/30

```

```

22/22 [=====] - 2s 78ms/step - loss: 2.2963 - accuracy: 0.1061 - val
Epoch 20/30
22/22 [=====] - 2s 79ms/step - loss: 2.2988 - accuracy: 0.1212 - val
Epoch 21/30
22/22 [=====] - 2s 74ms/step - loss: 2.2966 - accuracy: 0.1515 - val
Epoch 22/30
22/22 [=====] - 2s 79ms/step - loss: 2.2963 - accuracy: 0.0606 - val
Epoch 23/30
22/22 [=====] - 2s 77ms/step - loss: 2.2960 - accuracy: 0.1364 - val
Epoch 24/30
22/22 [=====] - 2s 78ms/step - loss: 2.2991 - accuracy: 0.1061 - val
Epoch 25/30
22/22 [=====] - 2s 75ms/step - loss: 2.2975 - accuracy: 0.1212 - val
Epoch 26/30
22/22 [=====] - 2s 78ms/step - loss: 2.2990 - accuracy: 0.0909 - val
Epoch 27/30
22/22 [=====] - 2s 75ms/step - loss: 2.2952 - accuracy: 0.1818 - val
Epoch 28/30
22/22 [=====] - 2s 78ms/step - loss: 2.2970 - accuracy: 0.1364 - val
Epoch 29/30

```

```
model.save('fruit_recognition.h5')
```

```
from keras.models import load_model
recognition = load_model('fruit_recognition.h5')
```

```
labels = {0: 'apple', 1: 'banana', 2: 'cherry', 3: 'coconut',
          4: 'durian', 5: 'kiwi', 6: 'mango', 7: 'orange',
          8: 'pomelo', 9: 'water melon'}
```

```
from keras.preprocessing.image import load_img, img_to_array
img = load_img('mango.jpg', target_size = (224,224))
plt.imshow(img)
img = img_to_array(img)
img = img.reshape(1,224,224,3)
img = img.astype('float32')
img = img/255
img.shape
```

```
val = recognition.predict(img)
np.argmax(val, axis=1)
print('This is ', labels[np.argmax(val)])
```

This is mango



